

Problem Statement: The telecom industry faces the challenge of customer churn, where customers cancel or do not renew their subscriptions, leading to potential revenue loss. Identifying potential churn customers is vital for the company's growth and customer retention.

Aim: The aim of this analysis is to develop predictive models to classify potential churn customers based on numerical and categorical features. The analysis involves binary classification, where customers are categorized as likely to churn (Yes) or not (No).

Importance of Churn Analysis: Churn analysis provides valuable insights into customer behavior and preferences. It helps companies build effective strategies to retain customers, enhance service quality, and conduct targeted marketing campaigns.

Strategy: The analysis focuses on building accurate predictive models that address class imbalance to make informed decisions about potential churn customers.

Benefits: The analysis empowers telecom companies to optimize customer retention efforts, reduce churn rates, and cultivate trust with customers.

Practical Use: The outcomes of churn analysis aid companies in developing customer-centric approaches, improving customer satisfaction, and ultimately driving business growth.

```
import os
import warnings
warnings.filterwarnings('ignore')
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

import os: This line imports the 'os' module, which provides a way of using operating system-dependent functionality.

import warnings: This line imports the 'warnings' module, allowing us to handle or suppress warning messages during the code execution.

warnings.filterwarnings('ignore'): This line sets a filter to ignore warning messages. It's a good practice to suppress warnings that might not be critical to our analysis.

import numpy as np: This line imports the 'numpy' library and assigns it the alias 'np'. 'numpy' is a powerful library for numerical computing and handling arrays and matrices.

import pandas as pd: This line imports the 'pandas' library and assigns it the alias 'pd'. 'pandas' is widely used for data manipulation and analysis, providing powerful data structures like DataFrames.

import matplotlib.pyplot as plt: This line imports the 'matplotlib.pyplot' module and assigns it the alias 'plt'. 'matplotlib' is a popular plotting library in Python, and 'pyplot' provides a MATLAB-like interface for creating visualizations.

import seaborn as sns: This line imports the 'seaborn' library and assigns it the alias 'sns'. 'seaborn' is another data visualization library built on top of 'matplotlib', offering additional high-level plotting functions and attractive styles.

```
dataset= pd.read_csv('/content/Telcom Data.csv')
```

Data Loading The above code reads the telecom dataset from a CSV (Comma Separated Values) file named "Telcom Data.csv" located in the '/content' directory. It uses the 'pd.read_csv()' function provided by the 'pandas' library to load the data and store it in a DataFrame called 'dataset'.

```
dataset.head()
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure
0	7590-VHVEG	Female	0	Yes	No	1
1	5575-GNVDE	Male	0	No	No	34
2	3668-QPYBK	Male	0	No	No	2
3	7795-CF0CW	Male	0	No	No	45
4	9237-HQITU	Female	0	No	No	2

	MultipleLines	InternetService	OnlineSecurity	...
0	No phone service	DSL	No	...
1	No	DSL	Yes	...
2	No	DSL	Yes	...
3	No phone service	DSL	Yes	...
4	No	Fiber optic	No	...

	TechSupport	StreamingTV	StreamingMovies	Contract
PaperlessBilling \				
0	No	No	No	Month-to-month
Yes				
1	No	No	No	One year
No				
2	No	No	No	Month-to-month
Yes				
3	Yes	No	No	One year
No				
4	No	No	No	Month-to-month
Yes				
	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	Electronic check	29.85	29.85	No
1	Mailed check	56.95	1889.5	No
2	Mailed check	53.85	108.15	Yes
3	Bank transfer (automatic)	42.30	1840.75	No
4	Electronic check	70.70	151.65	Yes
[5 rows x 21 columns]				

The dataset contains information about telecom customers and various attributes related to their interactions and services with the company. Here is a more detailed description of each attribute:

customerID: A unique identifier for each customer.

gender: Indicates whether the customer is male or female.

SeniorCitizen: A binary indicator (1 or 0) that specifies whether the customer is a senior citizen (1: Yes, 0: No).

Partner: A binary indicator (Yes or No) representing whether the customer has a partner.

Dependents: A binary indicator (Yes or No) showing whether the customer has dependents (e.g., family members).

tenure: The number of months the customer has stayed with the company.

PhoneService: A binary indicator (Yes or No) indicating whether the customer has a phone service.

MultipleLines: Indicates whether the customer has multiple phone lines (Yes, No, or No phone service).

InternetService: Specifies the customer's internet service provider (DSL, Fiber optic, or No internet service).

OnlineSecurity: A binary indicator (Yes or No) showing whether the customer has online security service or not (or No internet service).

OnlineBackup: A binary indicator (Yes or No) showing whether the customer has online backup service or not (or No internet service).

DeviceProtection: A binary indicator (Yes or No) showing whether the customer has device protection service or not (or No internet service).

TechSupport: A binary indicator (Yes or No) indicating whether the customer has tech support service or not (or No internet service).

StreamingTV: A binary indicator (Yes or No) showing whether the customer has streaming TV service or not (or No internet service).

StreamingMovies: A binary indicator (Yes or No) showing whether the customer has streaming movie service or not (or No internet service).

Contract: The contract term of the customer (Month-to-month, One year, Two years).

PaperlessBilling: A binary indicator (Yes or No) showing whether the customer has opted for paperless billing.

PaymentMethod: Specifies the customer's payment method (Electronic check, Mailed check, Bank transfer (automatic), Credit card (automatic)).

MonthlyCharges: The amount charged to the customer on a monthly basis.

TotalCharges: The total amount charged to the customer over their entire tenure.

Churn: The target variable that indicates whether the customer has churned (left the telecom service) or not (Yes or No).

The dataset is relevant for conducting customer churn analysis, as it contains valuable information about customer behavior and characteristics that can be used to predict whether a customer is likely to churn or not. By utilizing machine learning algorithms and techniques, telecom companies can gain insights into customer retention strategies and improve overall customer satisfaction.

```
dataset.shape
```

```
(7043, 21)
```

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 7043 entries, 0 to 7042
```

```
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	customerID	7043 non-null	object
1	gender	7043 non-null	object
2	SeniorCitizen	7043 non-null	int64
3	Partner	7043 non-null	object
4	Dependents	7043 non-null	object
5	tenure	7043 non-null	int64

```

6   PhoneService      7043 non-null object
7   MultipleLines     7043 non-null object
8   InternetService   7043 non-null object
9   OnlineSecurity     7043 non-null object
10  OnlineBackup       7043 non-null object
11  DeviceProtection   7043 non-null object
12  TechSupport        7043 non-null object
13  StreamingTV        7043 non-null object
14  StreamingMovies    7043 non-null object
15  Contract           7043 non-null object
16  PaperlessBilling   7043 non-null object
17  PaymentMethod      7043 non-null object
18  MonthlyCharges     7043 non-null float64
19  TotalCharges       7043 non-null object
20  Churn              7043 non-null object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB

```

The code `dataset.info()` provides information about the DataFrame 'dataset', which contains telecom customer data. It gives an overview of the columns in the DataFrame, the number of non-null values in each column, and the data types of each column.

The 'object' data type represents strings or text, while 'int64' and 'float64' represent integer and floating-point numeric values, respectively.

It is important to check the data types as they can influence data processing and analysis. For instance, 'TotalCharges' is expected to be a numeric column, but it is currently identified as 'object'. This suggests that there might be some non-numeric values in the 'TotalCharges' column. To handle this issue, the column may need to be converted to a numeric data type.

```
dataset['Churn'].value_counts()
```

```

No      5174
Yes     1869
Name: Churn, dtype: int64

```

The code `dataset['Churn'].value_counts()` is used to count the number of occurrences of each unique value in the 'Churn' column of the DataFrame 'dataset'. The 'Churn' column represents whether a customer churned (canceled or did not renew their subscription) or not.

'No': This value appears 5174 times in the 'Churn' column, indicating that there are 5174 customers who did not churn.

'Yes': This value appears 1869 times in the 'Churn' column, indicating that there are 1869 customers who churned. The `value_counts()` function is a convenient way to quickly understand the distribution of values in a categorical column like 'Churn'. This information is useful for gaining insights into the churn rate and understanding the proportion of customers who have churned compared to those who have not.

```
dataset['Churn'] = dataset['Churn'].replace({'Yes':1, 'No':0})
```

By using the `replace()` function with the argument `{'Yes': 1, 'No': 0}`, we are mapping 'Yes' to 1 and 'No' to 0, converting the column to a binary representation. This transformation is essential for binary classification tasks, as many machine learning algorithms require numeric input features and target labels.

```
dataset.columns

Index(['customerID', 'gender', 'SeniorCitizen', 'Partner',
      'Dependents',
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
      'TechSupport',
      'StreamingTV', 'StreamingMovies', 'Contract',
      'PaperlessBilling',
      'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')
```

```
for i in dataset.columns:
    print("*****", i,
          "*****")
    print()
    print(set(dataset[i].tolist()))
    print()
```

The code iterates over each column in the DataFrame 'dataset' and prints information about each column, including the column name, unique values present in the column, and the data type of the column.

The code provided is helpful because it allows you to quickly understand the unique values present in each column of the dataset. It provides insights into the nature of the data in each column, especially for categorical columns, which is crucial for data exploration and analysis. By printing the unique values for each column, you can gain a better understanding of the different categories or levels in the dataset.

Using AutoML Pycaret

PyCaret provides a high-level API that allows users to perform complex machine learning tasks with just a few lines of code. It automates many of the repetitive tasks involved in building and evaluating machine learning models, making it easier for data scientists and machine learning practitioners to experiment with different models and techniques.

```
!pip install pycaret

Collecting pycaret
  Downloading pycaret-3.0.4-py3-none-any.whl (484 kB)
  484.4/484.4 kB 8.3 MB/s eta
0:00:00
ent already satisfied: ipython>=5.5.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from pycaret) (7.34.0)
Requirement already satisfied: ipywidgets>=7.6.5 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (7.7.1)
Requirement already satisfied: tqdm>=4.62.0 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (4.65.0)
Requirement already satisfied: numpy<1.24,>=1.21 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (1.22.4)
Requirement already satisfied: pandas<2.0.0,>=1.3.0 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (1.5.3)
Requirement already satisfied: jinja2>=1.2 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (3.1.2)
Requirement already satisfied: scipy<2.0.0 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (1.10.1)
Requirement already satisfied: joblib>=1.2.0 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (1.3.1)
Requirement already satisfied: scikit-learn<1.3.0,>=1.0 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (1.2.2)
Collecting pyod>=1.0.8 (from pycaret)
  Downloading pyod-1.1.0.tar.gz (153 kB)
  153.4/153.4 kB 20.7 MB/s eta
0:00:00
etaddata (setup.py) ... ent already satisfied: imbalanced-learn>=0.8.1
in /usr/local/lib/python3.10/dist-packages (from pycaret) (0.10.1)
Collecting category-encoders>=2.4.0 (from pycaret)
  Downloading category_encoders-2.6.1-py2.py3-none-any.whl (81 kB)
  81.9/81.9 kB 10.8 MB/s eta
0:00:00
ent already satisfied: lightgbm>=3.0.0 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (3.3.5)
Requirement already satisfied: numba>=0.55.0 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (0.56.4)
Requirement already satisfied: requests>=2.27.1 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (2.27.1)
Requirement already satisfied: psutil>=5.9.0 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (5.9.5)
Requirement already satisfied: markupsafe>=2.0.1 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (2.1.3)
Collecting importlib-metadata>=4.12.0 (from pycaret)
  Downloading importlib_metadata-6.8.0-py3-none-any.whl (22 kB)
Requirement already satisfied: nbformat>=4.2.0 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (5.9.1)
Requirement already satisfied: cloudpickle in
/usr/local/lib/python3.10/dist-packages (from pycaret) (2.2.1)
Collecting deprecation>=2.1.0 (from pycaret)
  Downloading deprecation-2.1.0-py2.py3-none-any.whl (11 kB)
Collecting xxhash (from pycaret)
  Downloading xxhash-3.3.0-cp310-cp310-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (194 kB)
  194.1/194.1 kB 23.2 MB/s eta
```

```

0:00:00
ent already satisfied: matplotlib>=3.3.0 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (3.7.1)
Collecting scikit-plot>=0.3.7 (from pycaret)
  Downloading scikit_plot-0.3.7-py3-none-any.whl (33 kB)
Requirement already satisfied: yellowbrick>=1.4 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (1.5)
Requirement already satisfied: plotly>=5.0.0 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (5.13.1)
Collecting kaleido>=0.2.1 (from pycaret)
  Downloading kaleido-0.2.1-py2.py3-none-manylinux1_x86_64.whl (79.9
MB)
_____ 79.9/79.9 MB 11.1 MB/s eta
0:00:00
draw==0.15 (from pycaret)
  Downloading schemdraw-0.15-py3-none-any.whl (106 kB)
_____ 106.8/106.8 kB 14.8 MB/s eta
0:00:00
pler>=0.8.3.1 (from pycaret)
  Downloading plotly_resampler-0.9.1-py3-none-any.whl (73 kB)
_____ 73.4/73.4 kB 10.1 MB/s eta
0:00:00
ent already satisfied: statsmodels>=0.12.1 in
/usr/local/lib/python3.10/dist-packages (from pycaret) (0.13.5)
Collecting sktime!=0.17.1,!0.17.2,!0.18.0,>=0.16.1 (from pycaret)
  Downloading sktime-0.21.0-py3-none-any.whl (17.1 MB)
_____ 17.1/17.1 MB 52.1 MB/s eta
0:00:00
pycaret)
  Downloading tbats-1.1.3-py3-none-any.whl (44 kB)
_____ 44.0/44.0 kB 6.0 MB/s eta
0:00:00
darima!=1.8.1,<3.0.0,>=1.8.0 (from pycaret)
  Downloading pmdarima-2.0.3-cp310-cp310-
manylinux_2_17_x86_64.manylinux2014_x86_64.manylinux_2_28_x86_64.whl
(1.8 MB)
_____ 1.8/1.8 MB 56.8 MB/s eta
0:00:00
pycaret)
  Downloading wurlitzer-3.0.3-py3-none-any.whl (7.3 kB)
Requirement already satisfied: patsy>=0.5.1 in
/usr/local/lib/python3.10/dist-packages (from category-
encoders>=2.4.0->pycaret) (0.5.3)
Requirement already satisfied: packaging in
/usr/local/lib/python3.10/dist-packages (from deprecation>=2.1.0-
>pycaret) (23.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from imbalanced-learn>=0.8.1-
>pycaret) (3.2.0)

```



```
Requirement already satisfied: zipp>=0.5 in
/usr/local/lib/python3.10/dist-packages (from importlib-
metadata>=4.12.0->pycaret) (3.16.2)
Requirement already satisfied: setuptools>=18.5 in
/usr/local/lib/python3.10/dist-packages (from ipython>=5.5.0->pycaret)
(67.7.2)
Collecting jedi>=0.16 (from ipython>=5.5.0->pycaret)
  Downloading jedi-0.19.0-py2.py3-none-any.whl (1.6 MB)
_____ 1.6/1.6 MB 91.0 MB/s eta
0:00:00
Requirement already satisfied: decorator in /usr/local/lib/python3.10/dist-
packages (from ipython>=5.5.0->pycaret) (4.4.2)
Requirement already satisfied: pickleshare in
/usr/local/lib/python3.10/dist-packages (from ipython>=5.5.0->pycaret)
(0.7.5)
Requirement already satisfied: traitlets>=4.2 in
/usr/local/lib/python3.10/dist-packages (from ipython>=5.5.0->pycaret)
(5.7.1)
Requirement already satisfied: prompt-toolkit!=3.0.0,!
=3.0.1,<3.1.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from
ipython>=5.5.0->pycaret) (3.0.39)
Requirement already satisfied: pygments in
/usr/local/lib/python3.10/dist-packages (from ipython>=5.5.0->pycaret)
(2.14.0)
Requirement already satisfied: backcall in
/usr/local/lib/python3.10/dist-packages (from ipython>=5.5.0->pycaret)
(0.2.0)
Requirement already satisfied: matplotlib-inline in
/usr/local/lib/python3.10/dist-packages (from ipython>=5.5.0->pycaret)
(0.1.6)
Requirement already satisfied: pexpect>4.3 in
/usr/local/lib/python3.10/dist-packages (from ipython>=5.5.0->pycaret)
(4.8.0)
Requirement already satisfied: ipykernel>=4.5.1 in
/usr/local/lib/python3.10/dist-packages (from ipywidgets>=7.6.5-
>pycaret) (5.5.6)
Requirement already satisfied: ipython-genutils~=0.2.0 in
/usr/local/lib/python3.10/dist-packages (from ipywidgets>=7.6.5-
>pycaret) (0.2.0)
Requirement already satisfied: widgetsnbextension~=3.6.0 in
/usr/local/lib/python3.10/dist-packages (from ipywidgets>=7.6.5-
>pycaret) (3.6.4)
Requirement already satisfied: jupyterlab-widgets>=1.0.0 in
/usr/local/lib/python3.10/dist-packages (from ipywidgets>=7.6.5-
>pycaret) (3.0.8)
Requirement already satisfied: wheel in
/usr/local/lib/python3.10/dist-packages (from lightgbm>=3.0.0-
>pycaret) (0.41.0)
Requirement already satisfied: contourpy>=1.0.1 in
```

```

/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0-
>pycaret) (1.1.0)
Requirement already satisfied: cyclor>=0.10 in
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0-
>pycaret) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0-
>pycaret) (4.41.1)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0-
>pycaret) (1.4.4)
Requirement already satisfied: pillow>=6.2.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0-
>pycaret) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0-
>pycaret) (3.1.0)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0-
>pycaret) (2.8.2)
Requirement already satisfied: fastjsonschema in
/usr/local/lib/python3.10/dist-packages (from nbformat>=4.2.0-
>pycaret) (2.18.0)
Requirement already satisfied: jsonschema>=2.6 in
/usr/local/lib/python3.10/dist-packages (from nbformat>=4.2.0-
>pycaret) (4.3.3)
Requirement already satisfied: jupyter-core in
/usr/local/lib/python3.10/dist-packages (from nbformat>=4.2.0-
>pycaret) (5.3.1)
Requirement already satisfied: llvmlite<0.40,>=0.39.0dev0 in
/usr/local/lib/python3.10/dist-packages (from numba>=0.55.0->pycaret)
(0.39.1)
Requirement already satisfied: pytz>=2020.1 in
/usr/local/lib/python3.10/dist-packages (from pandas<2.0.0,>=1.3.0-
>pycaret) (2022.7.1)
Requirement already satisfied: tenacity>=6.2.0 in
/usr/local/lib/python3.10/dist-packages (from plotly>=5.0.0->pycaret)
(8.2.2)
Collecting dash<3.0.0,>=2.11.0 (from plotly-resampler>=0.8.3.1-
>pycaret)
  Downloading dash-2.11.1-py3-none-any.whl (10.4 MB)
  _____ 10.4/10.4 MB 105.3 MB/s eta
0:00:00
  plotly-resampler>=0.8.3.1->pycaret)
  Downloading orjson-3.9.2-cp310-cp310-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (138 kB)
  _____ 138.7/138.7 kB 16.3 MB/s eta
0:00:00
  plotly-resampler>=0.8.3.1->pycaret)

```

```

    Downloading trace_updater-0.0.9.1-py3-none-any.whl (185 kB)
    185.2/185.2 kB 20.8 MB/s eta
0:00:00
ple==0.1.2 (from plotly-resampler>=0.8.3.1->pycaret)
    Downloading tsdownsample-0.1.2-cp310-cp310-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (2.3 MB)
    2.3/2.3 MB 75.6 MB/s eta
0:00:00
ent already satisfied: Cython!=0.29.18,!>=0.29.31,>=0.29 in
/usr/local/lib/python3.10/dist-packages (from pmdarima!
=1.8.1,<3.0.0,>=1.8.0->pycaret) (0.29.36)
Requirement already satisfied: urllib3 in
/usr/local/lib/python3.10/dist-packages (from pmdarima!
=1.8.1,<3.0.0,>=1.8.0->pycaret) (1.26.16)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-
packages (from pyod>=1.0.8->pycaret) (1.16.0)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.27.1-
>pycaret) (2023.7.22)
Requirement already satisfied: charset-normalizer~>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.27.1-
>pycaret) (2.0.12)
Requirement already satisfied: idna<4,>=2.5 in
/usr/local/lib/python3.10/dist-packages (from requests>=2.27.1-
>pycaret) (3.4)
Collecting deprecated>=1.2.13 (from sktime!=0.17.1,!>=0.17.2,!
=0.18.0,>=0.16.1->pycaret)
    Downloading Deprecated-1.2.14-py2.py3-none-any.whl (9.6 kB)
Collecting scikit-base<0.6.0 (from sktime!=0.17.1,!>=0.17.2,!
=0.18.0,>=0.16.1->pycaret)
    Downloading scikit_base-0.5.0-py3-none-any.whl (118 kB)
    118.2/118.2 kB 13.7 MB/s eta
0:00:00
ent already satisfied: Flask<2.3.0,>=1.0.4 in
/usr/local/lib/python3.10/dist-packages (from dash<3.0.0,>=2.11.0-
>plotly-resampler>=0.8.3.1->pycaret) (2.2.5)
Collecting Werkzeug<2.3.0 (from dash<3.0.0,>=2.11.0->plotly-
resampler>=0.8.3.1->pycaret)
    Downloading Werkzeug-2.2.3-py3-none-any.whl (233 kB)
    233.6/233.6 kB 26.1 MB/s eta
0:00:00
l-components==2.0.0 (from dash<3.0.0,>=2.11.0->plotly-
resampler>=0.8.3.1->pycaret)
    Downloading dash_html_components-2.0.0-py3-none-any.whl (4.1 kB)
Collecting dash-core-components==2.0.0 (from dash<3.0.0,>=2.11.0-
>plotly-resampler>=0.8.3.1->pycaret)
    Downloading dash_core_components-2.0.0-py3-none-any.whl (3.8 kB)
Collecting dash-table==5.0.0 (from dash<3.0.0,>=2.11.0->plotly-
resampler>=0.8.3.1->pycaret)

```

Downloading dash_table-5.0.0-py3-none-any.whl (3.9 kB)
 Requirement already satisfied: typing-extensions>=4.1.1 in
 /usr/local/lib/python3.10/dist-packages (from dash<3.0.0,>=2.11.0-
 >plotly-resampler>=0.8.3.1->pycaret) (4.7.1)
 Collecting retrying (from dash<3.0.0,>=2.11.0->plotly-
 resampler>=0.8.3.1->pycaret)
 Downloading retrying-1.3.4-py3-none-any.whl (11 kB)
 Collecting ansi2html (from dash<3.0.0,>=2.11.0->plotly-
 resampler>=0.8.3.1->pycaret)
 Downloading ansi2html-1.8.0-py3-none-any.whl (16 kB)
 Requirement already satisfied: nest-asyncio in
 /usr/local/lib/python3.10/dist-packages (from dash<3.0.0,>=2.11.0-
 >plotly-resampler>=0.8.3.1->pycaret) (1.5.6)
 Requirement already satisfied: wrapt<2,>=1.10 in
 /usr/local/lib/python3.10/dist-packages (from deprecated>=1.2.13-
 >sktime!=0.17.1,!0.17.2,!0.18.0,>=0.16.1->pycaret) (1.14.1)
 Requirement already satisfied: jupyter-client in
 /usr/local/lib/python3.10/dist-packages (from ipykernel>=4.5.1-
 >ipywidgets>=7.6.5->pycaret) (6.1.12)
 Requirement already satisfied: tornado>=4.2 in
 /usr/local/lib/python3.10/dist-packages (from ipykernel>=4.5.1-
 >ipywidgets>=7.6.5->pycaret) (6.3.1)
 Requirement already satisfied: parso<0.9.0,>=0.8.3 in
 /usr/local/lib/python3.10/dist-packages (from jedi>=0.16-
 >ipython>=5.5.0->pycaret) (0.8.3)
 Requirement already satisfied: attrs>=17.4.0 in
 /usr/local/lib/python3.10/dist-packages (from jsonschema>=2.6-
 >nbformat>=4.2.0->pycaret) (23.1.0)
 Requirement already satisfied: pyparsing!=0.17.0,!0.17.1,!
 =0.17.2,>=0.14.0 in /usr/local/lib/python3.10/dist-packages (from
 jsonschema>=2.6->nbformat>=4.2.0->pycaret) (0.19.3)
 Requirement already satisfied: ptyprocess>=0.5 in
 /usr/local/lib/python3.10/dist-packages (from pexpect>4.3-
 >ipython>=5.5.0->pycaret) (0.7.0)
 Requirement already satisfied: wcwidth in
 /usr/local/lib/python3.10/dist-packages (from prompt-toolkit!=3.0.0,!
 =3.0.1,<3.1.0,>=2.0.0->ipython>=5.5.0->pycaret) (0.2.6)
 Requirement already satisfied: notebook>=4.4.1 in
 /usr/local/lib/python3.10/dist-packages (from
 widgetsnbextension~=3.6.0->ipywidgets>=7.6.5->pycaret) (6.4.8)
 Requirement already satisfied: platformdirs>=2.5 in
 /usr/local/lib/python3.10/dist-packages (from jupyter-core-
 >nbformat>=4.2.0->pycaret) (3.9.1)
 Requirement already satisfied: itsdangerous>=2.0 in
 /usr/local/lib/python3.10/dist-packages (from Flask<2.3.0,>=1.0.4-
 >dash<3.0.0,>=2.11.0->plotly-resampler>=0.8.3.1->pycaret) (2.1.2)
 Requirement already satisfied: click>=8.0 in
 /usr/local/lib/python3.10/dist-packages (from Flask<2.3.0,>=1.0.4-
 >dash<3.0.0,>=2.11.0->plotly-resampler>=0.8.3.1->pycaret) (8.1.6)

Requirement already satisfied: pyzmq>=17 in
/usr/local/lib/python3.10/dist-packages (from notebook>=4.4.1-
>widgetsnbextension~=3.6.0->ipywidgets>=7.6.5->pycaret) (23.2.1)

Requirement already satisfied: argon2-cffi in
/usr/local/lib/python3.10/dist-packages (from notebook>=4.4.1-
>widgetsnbextension~=3.6.0->ipywidgets>=7.6.5->pycaret) (21.3.0)

Requirement already satisfied: nbconvert in
/usr/local/lib/python3.10/dist-packages (from notebook>=4.4.1-
>widgetsnbextension~=3.6.0->ipywidgets>=7.6.5->pycaret) (6.5.4)

Requirement already satisfied: Send2Trash>=1.8.0 in
/usr/local/lib/python3.10/dist-packages (from notebook>=4.4.1-
>widgetsnbextension~=3.6.0->ipywidgets>=7.6.5->pycaret) (1.8.2)

Requirement already satisfied: terminado>=0.8.3 in
/usr/local/lib/python3.10/dist-packages (from notebook>=4.4.1-
>widgetsnbextension~=3.6.0->ipywidgets>=7.6.5->pycaret) (0.17.1)

Requirement already satisfied: prometheus-client in
/usr/local/lib/python3.10/dist-packages (from notebook>=4.4.1-
>widgetsnbextension~=3.6.0->ipywidgets>=7.6.5->pycaret) (0.17.1)

Requirement already satisfied: argon2-cffi-bindings in
/usr/local/lib/python3.10/dist-packages (from argon2-cffi-
>notebook>=4.4.1->widgetsnbextension~=3.6.0->ipywidgets>=7.6.5-
>pycaret) (21.2.0)

Requirement already satisfied: lxml in /usr/local/lib/python3.10/dist-
packages (from nbconvert->notebook>=4.4.1->widgetsnbextension~=3.6.0-
>ipywidgets>=7.6.5->pycaret) (4.9.3)

Requirement already satisfied: beautifulsoup4 in
/usr/local/lib/python3.10/dist-packages (from nbconvert-
>notebook>=4.4.1->widgetsnbextension~=3.6.0->ipywidgets>=7.6.5-
>pycaret) (4.11.2)

Requirement already satisfied: bleach in
/usr/local/lib/python3.10/dist-packages (from nbconvert-
>notebook>=4.4.1->widgetsnbextension~=3.6.0->ipywidgets>=7.6.5-
>pycaret) (6.0.0)

Requirement already satisfied: defusedxml in
/usr/local/lib/python3.10/dist-packages (from nbconvert-
>notebook>=4.4.1->widgetsnbextension~=3.6.0->ipywidgets>=7.6.5-
>pycaret) (0.7.1)

Requirement already satisfied: entrypoints>=0.2.2 in
/usr/local/lib/python3.10/dist-packages (from nbconvert-
>notebook>=4.4.1->widgetsnbextension~=3.6.0->ipywidgets>=7.6.5-
>pycaret) (0.4)

Requirement already satisfied: jupyterlab-pygments in
/usr/local/lib/python3.10/dist-packages (from nbconvert-
>notebook>=4.4.1->widgetsnbextension~=3.6.0->ipywidgets>=7.6.5-
>pycaret) (0.2.2)

Requirement already satisfied: mistune<2,>=0.8.1 in
/usr/local/lib/python3.10/dist-packages (from nbconvert-
>notebook>=4.4.1->widgetsnbextension~=3.6.0->ipywidgets>=7.6.5-
>pycaret) (0.8.4)

```

Requirement already satisfied: nbclient>=0.5.0 in
/usr/local/lib/python3.10/dist-packages (from nbconvert-
>notebook>=4.4.1->widgetsnbextension~=3.6.0->ipywidgets>=7.6.5-
>pycaret) (0.8.0)
Requirement already satisfied: pandocfilters>=1.4.1 in
/usr/local/lib/python3.10/dist-packages (from nbconvert-
>notebook>=4.4.1->widgetsnbextension~=3.6.0->ipywidgets>=7.6.5-
>pycaret) (1.5.0)
Requirement already satisfied: tinycss2 in
/usr/local/lib/python3.10/dist-packages (from nbconvert-
>notebook>=4.4.1->widgetsnbextension~=3.6.0->ipywidgets>=7.6.5-
>pycaret) (1.2.1)
Requirement already satisfied: cffi>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from argon2-cffi-bindings-
>argon2-cffi->notebook>=4.4.1->widgetsnbextension~=3.6.0-
>ipywidgets>=7.6.5->pycaret) (1.15.1)
Requirement already satisfied: soupsieve>1.2 in
/usr/local/lib/python3.10/dist-packages (from beautifulsoup4-
>nbconvert->notebook>=4.4.1->widgetsnbextension~=3.6.0-
>ipywidgets>=7.6.5->pycaret) (2.4.1)
Requirement already satisfied: webencodings in
/usr/local/lib/python3.10/dist-packages (from bleach->nbconvert-
>notebook>=4.4.1->widgetsnbextension~=3.6.0->ipywidgets>=7.6.5-
>pycaret) (0.5.1)
Requirement already satisfied: pyparser in
/usr/local/lib/python3.10/dist-packages (from cffi>=1.0.1->argon2-
cffi-bindings->argon2-cffi->notebook>=4.4.1-
>widgetsnbextension~=3.6.0->ipywidgets>=7.6.5->pycaret) (2.21)
Building wheels for collected packages: pyod
  Building wheel for pyod (setup.py) ... e=pyod-1.1.0-py3-none-any.whl
size=185330
sha256=f3bf183f051e4e73152a21c30338d5d277000bbc8207886662b7d03b8f4d8b7
a
  Stored in directory:
/root/.cache/pip/wheels/36/8e/e2/e932956b10b843eb6be9eefa70b5c1bee7b56
1be14c423b136
Successfully built pyod
Installing collected packages: trace-updater, kaleido, dash-table,
dash-html-components, dash-core-components, xxhash, wurlitzer,
Werkzeug, tsdownsample, scikit-base, schemdraw, retrying, orjson,
jedi, importlib-metadata, deprecation, deprecated, ansi2html, sktime,
scikit-plot, pyod, dash, pmdarima, plotly-resampler, category-
encoders, tbats, pycaret
  Attempting uninstall: Werkzeug
    Found existing installation: Werkzeug 2.3.6
    Uninstalling Werkzeug-2.3.6:
      Successfully uninstalled Werkzeug-2.3.6
  Attempting uninstall: importlib-metadata
    Found existing installation: importlib-metadata 4.6.4

```

```

Uninstalling importlib-metadata-4.6.4:
Successfully uninstalled importlib-metadata-4.6.4
Successfully installed Werkzeug-2.2.3 ansi2html-1.8.0 category-
encoders-2.6.1 dash-2.11.1 dash-core-components-2.0.0 dash-html-
components-2.0.0 dash-table-5.0.0 deprecated-1.2.14 deprecation-2.1.0
importlib-metadata-6.8.0 jedi-0.19.0 kaleido-0.2.1 orjson-3.9.2
plotly-resampler-0.9.1 pmdarima-2.0.3 pycaret-3.0.4 pyod-1.1.0
retrying-1.3.4 schemdraw-0.15 scikit-base-0.5.0 scikit-plot-0.3.7
sktime-0.21.0 tbats-1.1.3 trace-updater-0.0.9.1 tsdownsample-0.1.2
wurlitzer-3.0.3 xxhash-3.3.0

from pycaret.classification import *

exp_clf = setup(data=dataset, target= 'Churn', pca= True ,
pca_components=0.95 , session_id=123)

<pandas.io.formats.style.Styler at 0x7877b27193c0>

```

In the given code, we use the setup() function from PyCaret to set up the machine learning environment for the churn analysis problem using the telecom dataset (dataset). Here's what each parameter does:

data: The input dataset that contains the features and the target variable (Churn in this case).

target: The name of the target variable to be predicted, which is "Churn" in this dataset.

pca: A boolean parameter indicating whether Principal Component Analysis (PCA) should be applied to the data. In this case, it is set to True, which means PCA will be used to reduce the dimensionality of the features while retaining 95% of the variance.

pca_components: The desired percentage of variance to be preserved after PCA is applied. In this case, it is set to 0.95, meaning that the top principal components will be retained to represent 95% of the total variance in the data.

session_id: An optional parameter that sets a random seed for reproducibility. It is set to 123 in this case. By using setup(), PyCaret automatically performs various preprocessing steps, such as handling missing values, encoding categorical variables, splitting the data into training and testing sets, and applying PCA to reduce the feature dimensions. Additionally, it performs automatic feature selection and engineering based on the data and target variable.

After calling setup(), the machine learning environment is ready, and we can proceed with training and evaluating models using PyCaret's simple and efficient API.

```

#comparing models

compare_models()

<IPython.core.display.HTML object>

<pandas.io.formats.style.Styler at 0x78782f864dc0>

```

```
{"model_id": "235a24c2dd2243e08c4a24c5aff7f356", "version_major": 2, "version_minor": 0}
```

<IPython.core.display.HTML object>

```
GradientBoostingClassifier(ccp_alpha=0.0, criterion='friedman_mse',
init=None,
                           learning_rate=0.1, loss='log_loss',
max_depth=3,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0,
min_samples_leaf=1,
                           min_samples_split=2,
min_weight_fraction_leaf=0.0,
                           n_estimators=100, n_iter_no_change=None,
                           random_state=123, subsample=1.0,
tol=0.0001,
                           validation_fraction=0.1, verbose=0,
                           warm_start=False)
```

We use the `compare_models()` function from PyCaret to automatically train and evaluate multiple classification models on the dataset. This function helps us quickly identify the top-performing models based on default hyperparameters and performance metrics.

Here's what the `compare_models()` function does:

Training Multiple Models: The `compare_models()` function trains several classification models on the training data. The models used for comparison include popular algorithms such as Decision Tree, Random Forest, Logistic Regression, Gradient Boosting, K-Nearest Neighbors (KNN), Support Vector Machine (SVM), etc.

Cross-Validation: By default, PyCaret performs a stratified k-fold cross-validation (with k=10) during the model training process. Cross-validation helps to assess the model's performance on different subsets of the training data, reducing the risk of overfitting.

Evaluation Metrics: For each model, PyCaret calculates and displays various evaluation metrics, such as accuracy, area under the receiver operating characteristic curve (AUC-ROC), recall, precision, F1-score, kappa, and Matthews correlation coefficient (MCC). These metrics provide insights into the model's predictive performance.

Model Ranking: The `compare_models()` function ranks the models based on their performance on the given dataset. It displays the models in descending order of their average accuracy score by default.

Return Value: The function returns a table that shows the performance metrics of each trained model. The table helps us easily compare the performance of different models and select the most suitable one for further tuning.

Using `compare_models()` is a powerful and time-saving approach to quickly identify promising models for the given classification problem. After analyzing the results, we can choose the best-performing model and proceed with further tuning and evaluation to optimize its performance.


```

train_data = dataset.sample(frac=0.8,
random_state=101).reset_index(drop=True)
test_data = dataset.drop(train_data.index).reset_index(drop=True)

#since gbc is giving best results we are proceeding with gbc i.e
gradient boosting classifier
gbc = create_model('gbc')

<IPython.core.display.HTML object>

<pandas.io.formats.style.Styler at 0x787828245a20>

{"model_id":"d66d6eb22ab348a1b9abb074df230a60","version_major":2,"version_minor":0}

<IPython.core.display.HTML object>

```

we use the create_model() function from PyCaret to create a Gradient Boosting Classifier (GBC) model. This function automatically trains the GBC model on the training data, performs cross-validation, and returns the trained model object.

```

tuned_gbc = tune_model(gbc)

<IPython.core.display.HTML object>

<pandas.io.formats.style.Styler at 0x787828246fb0>

{"model_id":"e24b0a32c0ce4443bb7ee9a90d2c204f","version_major":2,"version_minor":0}

Fitting 10 folds for each of 10 candidates, totalling 100 fits

<IPython.core.display.HTML object>

Original model was better than the tuned model, hence it will be
returned. NOTE: The display metrics are for the tuned model (not the
original one).

```

In the code tuned_gbc = tune_model(gbc), we are using the tune_model() function from PyCaret to perform hyperparameter tuning on the Gradient Boosting Classifier (gbc) model that we created earlier.

Hyperparameter tuning is the process of finding the best set of hyperparameters for a machine learning model that results in optimal performance. Hyperparameters are model settings that are not learned during training and need to be specified before training the model. By tuning the hyperparameters, we can improve the model's performance and make it more suitable for our specific dataset.

When we call tune_model(gbc), PyCaret will automatically search for the best hyperparameters for the Gradient Boosting Classifier using a technique called "Random Grid Search." It will create a grid of possible hyperparameter values and randomly sample combinations from this grid to

train and evaluate the model. The combination of hyperparameters that gives the best performance will be returned as the `tuned_gbc` model.

By tuning the hyperparameters, we can potentially improve the accuracy and generalization of the Gradient Boosting Classifier model on our specific dataset. After hyperparameter tuning, the `tuned_gbc` model will be ready for evaluation and prediction.

We can now use the `evaluate_model()` function to assess the performance of the tuned model on the training data and use the `predict_model()` function to make predictions on new, unseen data. This helps us understand how well the model is generalizing to unseen data and how it performs in real-world scenarios.

When we used the `tune_model(gbc)` function, PyCaret performed hyperparameter tuning on the `gbc` model and evaluated its performance using cross-validation and evaluation metrics. However, it found that the original `gbc` model, without hyperparameter tuning, had better performance compared to the tuned model.

Overall, it's not uncommon to encounter cases where hyperparameter tuning does not lead to significant improvements in model performance. In such cases, sticking with the original model is a reasonable choice to avoid overfitting and maintain simplicity.

```
#sticking to original model for evaluation  
evaluate_model(gbc)  
  
{"model_id":"6e84fa63b4374641aac66139e4766128","version_major":2,"version_minor":0}
```

we are using the `evaluate_model()` function from PyCaret to assess the performance of the Gradient Boosting Classifier (`gbc`) model on the training data. This function provides a comprehensive report containing various evaluation metrics that help us understand how well the model is performing.

When we call `evaluate_model(gbc)`, PyCaret performs the following tasks:

Computes Accuracy: The accuracy metric measures the proportion of correctly classified instances out of the total instances. It tells us how often the model makes correct predictions.

Computes AUC (Area Under the Curve): AUC is a performance metric for binary classification models that measures the area under the Receiver Operating Characteristic (ROC) curve. It provides an overall measure of the model's ability to discriminate between the positive and negative classes.

Computes Recall: Recall, also known as sensitivity or true positive rate, measures the proportion of actual positive instances correctly predicted by the model. It indicates how well the model identifies positive cases.

Computes Precision: Precision is the proportion of true positive instances out of the total instances predicted as positive by the model. It measures the accuracy of positive predictions made by the model.

Computes F1 Score: The F1 score is the harmonic mean of precision and recall. It provides a balance between precision and recall, especially in imbalanced datasets.

Computes Kappa: The Kappa statistic measures the agreement between the actual and predicted classifications, taking into account the possibility of agreement occurring by chance.

Computes MCC (Matthews Correlation Coefficient): MCC is a measure of the quality of binary classifications. It takes into account true and false positives and negatives and is considered a balanced metric for imbalanced datasets.

By evaluating the gbc model using `evaluate_model(gbc)`, we can get a comprehensive understanding of how well the model is performing and how it is handling the classification task on the training data. This analysis helps us identify potential areas for improvement or model adjustments to optimize its performance on unseen data.

Running the code in Google Colab can be advantageous as it provides faster execution times compared to local Jupyter notebooks due to the availability of better hardware resources and integration with Google Cloud services. Additionally, Colab allows us to leverage GPUs and TPUs, which can significantly speed up computations for machine learning tasks.

In the given code, we are using the `evaluate_model()` function from PyCaret to assess the performance of the Gradient Boosting Classifier (gbc) model on the training data. This function provides a comprehensive report containing various evaluation metrics that help us understand how well the model is performing.

```
predict_model(gbc)
```

```
<pandas.io.formats.style.Styler at 0x7877b278d6f0>
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	\
2937	5312-IRCFR	Female	0	Yes	Yes	64	
3276	4210-QFJMF	Female	0	No	No	4	
4374	2718-YSKCS	Male	0	Yes	Yes	71	
4375	9896-UYMIE	Male	0	No	No	66	
237	9903-LYSAB	Male	0	Yes	No	18	
...
4496	9489-JMTTN	Female	0	Yes	Yes	72	
921	8942-DBMHZ	Male	0	No	No	12	
5904	9402-CXWPL	Female	0	No	No	70	
3088	5751-USDBL	Male	0	Yes	Yes	46	
253	3282-ZISZV	Male	0	No	Yes	32	

	PhoneService	MultipleLines	InternetService
OnlineSecurity	...	\	
2937	Yes	Yes	Fiber optic
No	...		
3276	Yes	No	Fiber optic
No	...		
4374	Yes	No	No
service	...		No internet
4375	Yes	Yes	Fiber optic

Yes	...			
237		Yes	Yes	Fiber optic
No	...			
...	
...				
4496		Yes	Yes	DSL
Yes	...			
921		Yes	No	No No internet service
...	...			
5904		Yes	Yes	Fiber optic
No	...			
3088		Yes	No	DSL
Yes	...			
253		Yes	Yes	DSL
No	...			

		StreamingTV	StreamingMovies	Contract	\
2937		No	Yes	One year	
3276		No	Yes	Month-to-month	
4374	No internet service	No internet service		Two year	
4375		Yes	Yes	One year	
237		No	No	Month-to-month	
...		
4496		Yes	Yes	Two year	
921	No internet service	No internet service		Month-to-month	
5904		Yes	Yes	One year	
3088		Yes	Yes	Two year	
253		Yes	Yes	One year	

	PaperlessBilling	PaymentMethod	MonthlyCharges
TotalCharges	\		
2937	Yes	Electronic check	92.849998
5980.75			
3276	Yes	Electronic check	79.150002
317.25			
4374	Yes	Bank transfer (automatic)	19.600000
1387.45			
4375	Yes	Bank transfer (automatic)	114.300003
7383.7			
237	Yes	Electronic check	73.150002
1305.95			
...
...			
4496	No	Credit card (automatic)	89.750000
6595.9			
921	No	Mailed check	20.450001
255.35			
5904	No	Electronic check	98.900002
6838.6			

3088	No	Mailed check	81.000000
3846.35			
253	No	Credit card (automatic)	83.699997
2633.3			

	Churn	prediction_label	prediction_score
2937	0	0	0.9048
3276	1	1	0.6998
4374	0	0	0.9776
4375	0	0	0.9390
237	0	0	0.6802
...
4496	0	0	0.9048
921	0	0	0.9074
5904	0	0	0.8913
3088	0	0	0.7811
253	0	0	0.7097

[2113 rows x 23 columns]

By using `predict_model(gbc)`, we can gain insights into how well the trained model performs on new, unseen data. This allows us to understand its generalization capabilities and how it would behave in real-world scenarios. The predictions can be further analyzed and used for making business decisions, such as identifying potential churn customers and implementing targeted retention strategies to reduce churn rate.

```
unseen_prediction = predict_model(gbc , data = test_data)
unseen_prediction
```

<pandas.io.formats.style.Styler at 0x7877b278f670>

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	\
0	2320-JRSDE	Female	0	Yes	Yes	1	
1	2087-QAREY	Female	0	Yes	No	22	
2	0601-WZHJF	Male	0	Yes	No	14	
3	4423-JWZJN	Male	0	Yes	Yes	64	
4	5143-WMWO	Male	0	No	No	1	
...	
1404	6840-RESVB	Male	0	Yes	Yes	24	
1405	2234-XADUH	Female	0	Yes	Yes	72	
1406	4801-JZAZL	Female	0	Yes	Yes	11	
1407	8361-LTMKD	Male	1	Yes	No	4	
1408	3186-AJIEK	Male	0	No	No	66	

	PhoneService	MultipleLines	InternetService
OnlineSecurity	...	\	
0	Yes	No	No internet
service	...		
1	Yes	No	DSL
No	...		

2	No	No phone service	DSL
No ...			
3	Yes	Yes	Fiber optic
No ...			
4	Yes	No	No No internet service
...
1404	Yes	Yes	DSL
Yes ...			
1405	Yes	Yes	Fiber optic
No ...			
1406	No	No phone service	DSL
Yes ...			
1407	Yes	Yes	Fiber optic
No ...			
1408	Yes	No	Fiber optic
Yes ...			
	StreamingTV	StreamingMovies	Contract \
0	No internet service	No internet service	Month-to-month
1	No	No	Month-to-month
2	Yes	Yes	Month-to-month
3	No	Yes	One year
4	No internet service	No internet service	Month-to-month
...
1404	Yes	Yes	One year
1405	Yes	Yes	One year
1406	No	No	Month-to-month
1407	No	No	Month-to-month
1408	Yes	Yes	Two year
	PaperlessBilling	PaymentMethod	MonthlyCharges
TotalCharges \			
0	Yes	Electronic check	19.900000
19.9			
1	Yes	Mailed check	54.700001
1178.75			
2	No	Electronic check	46.349998
667.7			
3	No	Credit card (automatic)	90.250000
5629.15			
4	No	Electronic check	19.950001
19.95			
...
...			
1404	Yes	Mailed check	84.800003
1990.5			
1405	Yes	Credit card (automatic)	103.199997

```

7362.9
1406          Yes          Electronic check      29.600000
346.45
1407          Yes          Mailed check         74.400002
306.6
1408          Yes  Bank transfer (automatic)    105.650002
6844.5

```

```

      Churn  prediction_label  prediction_score
0         1                0             0.5923
1         0                0             0.8096
2         1                0             0.7387
3         0                0             0.9407
4         1                0             0.6575
...      ...                ...              ...
1404       0                0             0.6241
1405       0                0             0.8876
1406       0                0             0.8415
1407       1                1             0.5972
1408       0                0             0.8009

```

```
[1409 rows x 23 columns]
```

```

unseen_prediction1 = predict_model(gbc , data = train_data)
unseen_prediction1

```

```
<pandas.io.formats.style.Styler at 0x787828167d30>
```

```

      customerID  gender  SeniorCitizen  Partner  Dependents  tenure  \
0      8659-I00PU  Female                0      Yes          Yes      71
1      0887-HJGAR   Male                0      No           No       1
2      1029-QFBEN   Male                0      No           No       1
3      7579-00PEC  Female                1      Yes          No       2
4      8473-VUVJN   Male                1      No           No       1
...      ...      ...      ...      ...      ...      ...
5629  1428-GTBJJ   Male                0      No           No      11
5630  7310-EGVHZ   Male                0      No           No       1
5631  3045-XETSH  Female                0      No           No      10
5632  2237-ZFSMY  Female                0      No           No      39
5633  0178-CIIKR  Female                0      No           No       3

```

```

      PhoneService  MultipleLines  InternetService
OnlineSecurity    ...  \
0          Yes          Yes      Fiber optic
Yes  ...
1          Yes          No          DSL
No  ...
2          Yes          No          No  No internet
service  ...
3          Yes          No          DSL

```

No	...				
4		Yes	Yes	Fiber optic	
No	...				
...	
...					
5629		Yes	No	Fiber optic	
No	...				
5630		Yes	No	No	No internet service
5631		Yes	Yes	Fiber optic	
No	...				
5632		Yes	No	Fiber optic	
Yes	...				
5633		Yes	No	No	No internet service
...					

	StreamingTV	StreamingMovies	Contract \
0	No	Yes	Two year
1	No	No	Month-to-month
2	No internet service	No internet service	Month-to-month
3	No	No	Month-to-month
4	No	No	Month-to-month
...
5629	No	No	Month-to-month
5630	No internet service	No internet service	Month-to-month
5631	No	Yes	Month-to-month
5632	Yes	No	One year
5633	No internet service	No internet service	Month-to-month

	PaperlessBilling	PaymentMethod	MonthlyCharges
TotalCharges \			
0	No	Electronic check	100.449997
7159.7			
1	Yes	Mailed check	45.700001
45.7			
2	No	Mailed check	19.549999
19.55			
3	No	Credit card (automatic)	50.150002
115.1			
4	Yes	Electronic check	73.650002
73.65			
...
...			
5629	Yes	Electronic check	74.550003
824.75			
5630	No	Bank transfer (automatic)	20.200001
20.2			
5631	Yes	Electronic check	94.849998
953.45			

5632	Yes	Electronic check	95.550003
3692.85			
5633	No	Mailed check	19.950001
58			

	Churn	prediction_label	prediction_score
0	0	0	0.8913
1	1	1	0.5702
2	1	0	0.6972
3	1	1	0.5598
4	1	1	0.6372
...
5629	1	0	0.5680
5630	0	1	0.5234
5631	1	1	0.7292
5632	1	0	0.6720
5633	0	0	0.7939

```
[5634 rows x 23 columns]
```

By using `predict_model(gbc, data=test_data)`, we can now analyze and interpret the model's performance on new, unseen data. This allows us to understand how well the model generalizes to new customers and how accurately it predicts churn behavior on real-world data.

It is essential to evaluate the predictions on unseen data and compare them with the actual outcomes to assess the model's effectiveness and reliability. This validation process helps in identifying potential areas for improvement and fine-tuning the model to enhance its performance for future predictions.

```
save_model(gbc , 'gbc model')
```

Transformation Pipeline and Model Successfully Saved

[illegible]

```
transformer=SimpleImputer(add_indicator=False,
```

```
copy=True,
```

```
fill value=None,
```

```
keep_empty_features=False,
```

```
missing_values=nan,
```

```

strategy='mean',
verbose='deprecated'))),
    ('categorical_imputer',
     Transform...
     criterion='friedman_mse',
init=None,
     learning_rate=0.1,
loss='log_loss',
     max_depth=3,
max_features=None,
     max_leaf_nodes=None,
min_impurity_decrease=0.0,
     min_samples_leaf=1,
     min_samples_split=2,
min_weight_fraction_leaf=0.0,
     n_estimators=100,
     n_iter_no_change=None,
     random_state=123,
subsample=1.0,
     tol=0.0001,
validation_fraction=0.1,
     verbose=0,
warm_start=False)]],
    verbose=False),
    'gbc_model.pkl')

```